

```
In [ ]: import Pkg
Pkg.activate(@__DIR__)
Pkg.instantiate()
using LinearAlgebra, Plots
import ForwardDiff as FD
import MeshCat as mc
using Test
using Random
import Convex as cvx
import ECOS      # the solver we use in this hw
# import Hypatia # other solvers you can try
# import COSMO  # other solvers you can try
using ProgressMeter
include(joinpath(@__DIR__, "utils/rendezvous.jl"))
```

```
Activating project at `~/Desktop/2024Spring/CMU16745_OptimalControl/CMU16-745-Optimal-Control-HW/hw2`
Precompiling project...
✓ StaticArrayInterface
✓ CloseOpenIntervals
✓ StaticArrayInterface → StaticArrayInterfaceOffsetArraysExt
✓ LayoutPointers
✓ StaticArrayInterface → StaticArrayInterfaceStaticArraysExt
✓ SparseDiffTools
✓ StrideArraysCore
✓ Polyester
✓ FastBroadcast
✓ SparseDiffTools → SparseDiffToolsPolyesterExt
✓ DiffEqBase
✓ VectorizationBase
✓ SciMLNLSolve
✓ DiffEqCallbacks
✓ DiffEqBase → DiffEqBaseUnitfulExt
✓ SLEEPPirates
✓ ControlSystemsBase
✓ LoopVectorization
✓ SimpleNonlinearSolve
✓ LoopVectorization → SpecialFunctionsExt
✓ TriangularSolve
✓ SimpleNonlinearSolve → SimpleNonlinearSolveStaticArraysExt
✓ RecursiveFactorization
✓ LinearSolve
✓ LinearSolve → LinearSolveRecursiveArrayToolsExt
✓ LinearSolve → LinearSolveIterativeSolversExt
✓ NonlinearSolve
✓ NonlinearSolve → NonlinearSolveNLSolveExt
✓ OrdinaryDiffEq
✓ DelayDiffEq
✓ ControlSystems
31 dependencies successfully precompiled in 185 seconds. 302 already precompiled.
[ Info: Precompiling IJuliaExt [2f4121a4-3b3a-5ce6-9c5e-1f2673ce168a]
```

```
Out[ ]: thruster_model (generic function with 1 method)
```

Notes:

1. Some of the cells below will have multiple outputs (plots and animations), it can be easier to see everything if you do **Cell -> All Output -> Toggle Scrolling**, so that it simply expands the output area to match the size of the outputs.
2. Things in space move very slowly (by design), because of this, you may want to speed up the animations when you're viewing them. You can do this in MeshCat by doing **Open Controls -> Animations -> Time Scale**, to modify the time scale. You can also play/pause/scrub from this menu as well.
3. You can move around your view in MeshCat by **clicking + dragging**, and you can pan with **right click + dragging**, and zoom with the scroll wheel on your mouse (or trackpad specific alternatives).

```
In [ ]: # utilities for converting to and from vector of vectors <-> matrix
function mat_from_vec(X::Vector{Vector{Float64}})::Matrix
    # convert a vector of vectors to a matrix
    Xm = hcat(X...)
    return Xm
end
function vec_from_mat(Xm::Matrix)::Vector{Vector{Float64}}
    # convert a matrix into a vector of vectors
    X = [Xm[:,i] for i = 1:size(Xm,2)]
    return X
end
```

```
Out [ ]: vec_from_mat (generic function with 1 method)
```

Is LQR the answer for everything?

Unfortunately, no. LQR is great for problems with true quadratic costs and linear dynamics, but this is a very small subset of convex trajectory optimization problems. While a quadratic cost is common in control, there are other available convex cost functions that may better motivate the desired behavior of the system. These costs can be things like an L1 norm on the control inputs ($\|u\|_1$), or an L2 goal error ($\|x - x_{goal}\|_2$). Also, control problems often have constraints like path constraints, control bounds, or terminal constraints, that can't be handled with LQR. With the addition of these constraints, the trajectory optimization problem is still convex and easy to solve, but we can no longer just get an optimal gain K and apply a feedback policy in these situations.

The solution to this is Model Predictive Control (MPC). In MPC, we are setting up and solving a convex trajectory optimization at every time step, optimizing over some horizon or window into the future, and executing the first control in the solution. To see how this works, we are going to try this for a classic space control problem: the rendezvous.

Q3: Optimal Rendezvous and Docking (55 pts)

In this example, we are going to use convex optimization to control the SpaceX Dragon 1 spacecraft as it docks with the International Space Station (ISS). The dynamics of the Dragon vehicle can be modeled with [Clohessy-Wiltshire equations](#), which is a linear dynamics model in continuous time. The state and control of this system are the following:

$$x = [r_x, r_y, r_z, v_x, v_y, v_z]^T, \quad (1)$$

$$u = [t_x, t_y, t_z]^T, \quad (2)$$

where r is a relative position of the Dragon spacecraft with respect to the ISS, v is the relative velocity, and t is the thrust on the spacecraft. The continuous time dynamics of the vehicle are the following:

$$\dot{x} = \begin{bmatrix} 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \\ 3n^2 & 0 & 0 & 0 & 2n & 0 \\ 0 & 0 & 0 & -2n & 0 & 0 \\ 0 & 0 & -n^2 & 0 & 0 & 0 \end{bmatrix} x + \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} u, \quad (3)$$

where $n = \sqrt{\mu/a^3}$, with μ being the [standard gravitational parameter](#), and a being the semi-major axis of the orbit of the ISS.

We are going to use three different techniques for solving this control problem, the first is LQR, the second is convex trajectory optimization, and the third is convex MPC where we will be able to account for unmodeled dynamics in our system (the "sim to real" gap).

Part A: Discretize the dynamics (5 pts)

Use the matrix exponential to convert the linear ODE into a linear discrete time model (hint: the matrix exponential is just `exp()` in Julia when called on a matrix.

```
In [ ]: function create_dynamics(dt::Real)::Tuple{Matrix,Matrix}
    mu = 3.986004418e14 # standard gravitational parameter
    a = 6971100.0      # semi-major axis of ISS
    n = sqrt(mu/a^3)   # mean motion

    # continuous time dynamics  $\dot{x} = Ax + Bu$ 
    A = [0      0  0      1  0  0;
          0      0  0      0  1  0;
          0      0  0      0  0  1;
          3*n^2  0  0      0  2*n 0;
          0      0  0     -2*n 0  0;
          0      0 -n^2  0  0  0]

    B = Matrix([zeros(3,3);0.1*I(3)])

    # TODO: convert to discrete time  $X_{k+1} = A_d x_k + B_d u_k$ 
    Z = [A B; zeros(3,6) zeros(3,3)]
    Zdt = exp(Z*dt)
    A_d = Zdt[1:6,1:6]
    B_d = Zdt[1:6,7:9]

    return A_d, B_d
end
```

```
Out[ ]: create_dynamics (generic function with 1 method)
```

```
In [ ]: @testset "discrete dynamics" begin
    A,B = create_dynamics(1.0)

    x = [1,3,-.3,.2,.4,-.5]
    u = [-.1,.5,.3]

    # test these matrices
    @test isapprox(A*x + B*u, [1.195453, 3.424786, -0.78499972, 0.190925, 0.4495759, -0.4699993], atol = 1e-3)
    @test isapprox(det(A), 1, atol = 1e-8)
```

```
@test isapprox(norm(B,Inf), 0.0999999803, atol = 1e-5)
```

```
end
```

```
Test Summary: | Pass Total Time
discrete dynamics | 3 3 1.8s
```

```
Out [ ]: Test.DefaultTestSet("discrete dynamics", Any[], 3, false, false, true, 1.708630128535368e9, 1.708630130369195e9, false)
```

Part B: LQR (10 pts)

Now we will take a given reference trajectory `X_ref` and track it with finite-horizon LQR. Remember that finite-horizon LQR is solving this problem:

$$\min_{x_{1:N}, u_{1:N-1}} \sum_{i=1}^{N-1} \left[\frac{1}{2} (x_i - x_{ref,i})^T Q (x_i - x_{ref,i}) + \frac{1}{2} u_i^T R u_i \right] + \frac{1}{2} (x_N - x_{ref,N})^T Q_f (x_N - x_{ref,N}) \quad (4)$$

$$\text{st } x_1 = x_{IC} \quad (5)$$

$$x_{i+1} = Ax_i + Bu_i \quad \text{for } i = 1, 2, \dots, N-1 \quad (6)$$

where our policy is $u_i = -K_i(x_i - x_{ref,i})$. Use your code from the previous problem with your `fh_lqr` function to generate your gain matrices.

One twist we will throw into this is control constraints `u_min` and `u_max`. You should use the function `clamp.(u, u_min, u_max)` to clamp the values of your `u` to be within this range.

If implemented correctly, you should see the Dragon spacecraft dock with the ISS successfully, but only after it crashes through the ISS a little bit.

```
In [ ]: function fh_lqr(A::Matrix, # A matrix
    B::Matrix, # B matrix
    Q::Matrix, # cost weight
    R::Matrix, # cost weight
    Qf::Matrix, # term cost weight
    N::Int64   # horizon size
```

```

)::Tuple{Vector{Matrix{Float64}}, Vector{Matrix{Float64}}} # return two matrices

# check sizes of everything
nx,nu = size(B)
@assert size(A) == (nx, nx)
@assert size(Q) == (nx, nx)
@assert size(R) == (nu, nu)
@assert size(Qf) == (nx, nx)

# instantiate S and K
P = [zeros(nx,nx) for i = 1:N]
K = [zeros(nu,nx) for i = 1:N-1]

# initialize S[N] with Qf
P[N] = deepcopy(Qf)

# Ricatti
for k = N-1:-1:1
P[k] = Q + A'*P[k+1]*A - A'*P[k+1]*B*inv(R + B'*P[k+1]*B)*B'*P[k+1]*A
K[k] = inv(R + B'*P[k+1]*B)*B'*P[k+1]*A
end

return P, K
end

```

Out[]: fhlqr (generic function with 1 method)

In []: @testset "LQR rendezvous" begin

```

# create our discrete time model
dt = 1.0
A,B = create_dynamics(dt)

# get our sizes for state and control
nx,nu = size(B)

# initial and goal states
x0 = [-2;-4;2;0;0;.0]
xg = [0,-.68,3.05,0,0,0]

```

```

# bounds on U
u_max = 0.4
u_min = -u_max

# problem size and reference trajectory
N = 120
t_vec = 0:dt:((N-1)*dt)
X_ref = desired_trajectory_long(x0,xg,200,dt)[1:N]

# TODO: FHLQR
Q = diagm(ones(nx))
R = diagm(ones(nu))
Qf = 10*Q
# TODO get K's from fhlqr
Ps, Ks = fhlqr(A,B,Q,R,Qf,N)

# simulation
X_sim = [zeros(nx) for i = 1:N]
U_sim = [zeros(nu) for i = 1:N-1]
X_sim[1] = x0
for i = 1:(N-1)
    # TODO: put LQR control law here
    # make sure to clamp
    U_sim[i] = clamp.(-Ks[i]*(X_sim[i] - X_ref[i]),u_min,u_max)

    # simulate 1 step
    X_sim[i+1] = A*X_sim[i] + B*U_sim[i]
end

# -----plotting/animation-----
Xm = mat_from_vec(X_sim)
Um = mat_from_vec(U_sim)
display(plot(t_vec,Xm[1:3,:]',title = "Positions (LQR)",
            xlabel = "time (s)", ylabel = "position (m)",
            label = ["x" "y" "z"]))
display(plot(t_vec,Xm[4:6,:]',title = "Velocities (LQR)",
            xlabel = "time (s)", ylabel = "velocity (m/s)",

```



```

        label = ["x" "y" "z"]))
display(plot(t_vec[1:end-1],Um',title = "Control (LQR)",
        xlabel = "time (s)", ylabel = "thrust (N)",
        label = ["x" "y" "z"]))

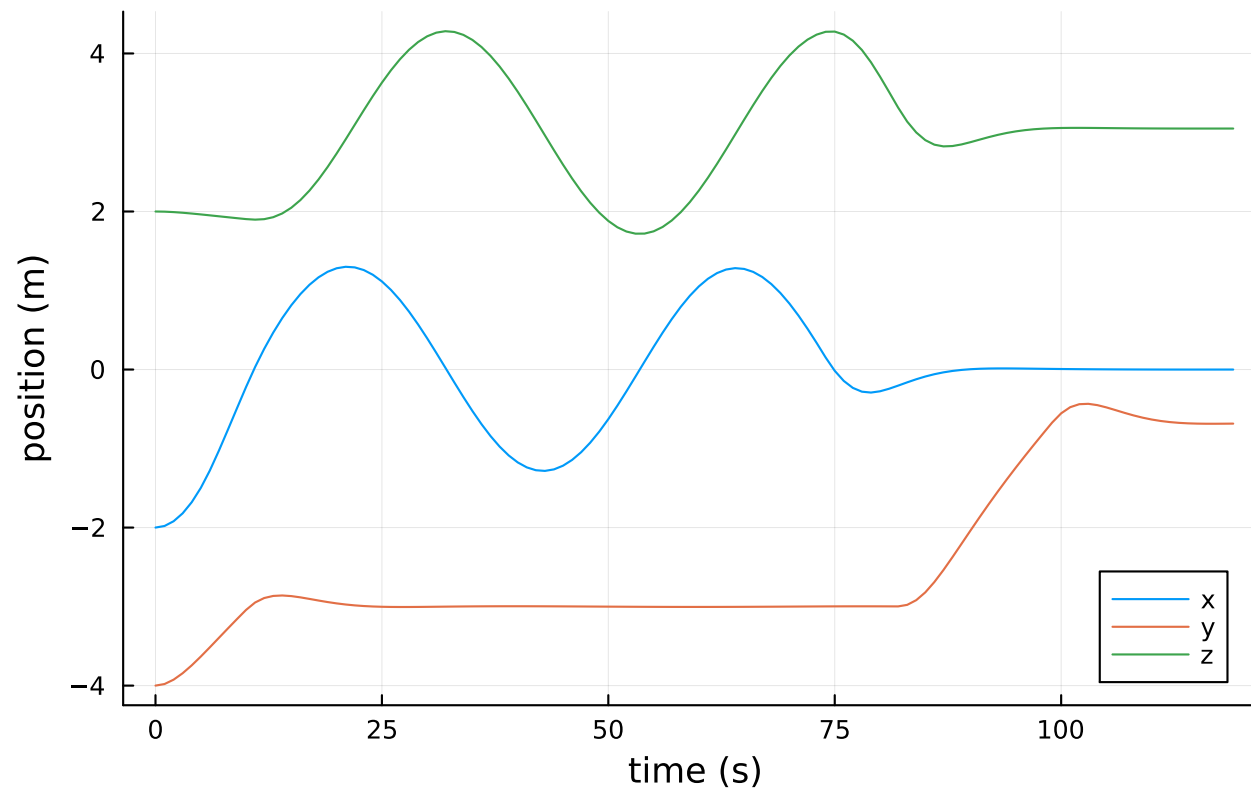
# feel free to toggle `show_reference`
display(animate_rendezvous(X_sim, X_ref, dt;show_reference = false))
# -----plotting/animation-----

# testing
xs=[x[1] for x in X_sim]
ys=[x[2] for x in X_sim]
zs=[x[3] for x in X_sim]
@test norm(X_sim[end] - xg) < .01 # goal
@test (xg[2] + .1) < maximum(ys) < 0 # we should have hit the ISS
@test maximum(zs) >= 4 # check to see if you did the circle
@test minimum(zs) <= 2 # check to see if you did the circle
@test maximum(xs) >= 1 # check to see if you did the circle
@test maximum(norm.(U_sim,Inf)) <= 0.4 # control constraints satisfied

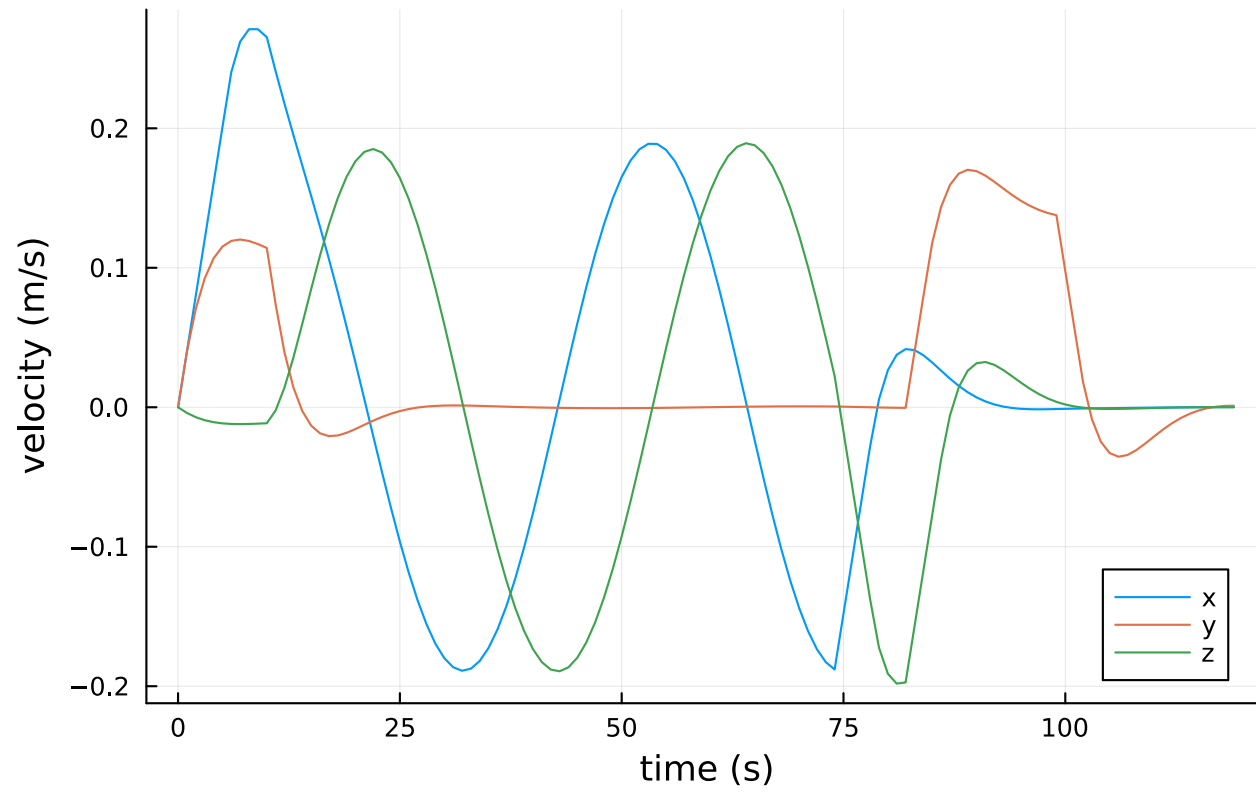
```

end

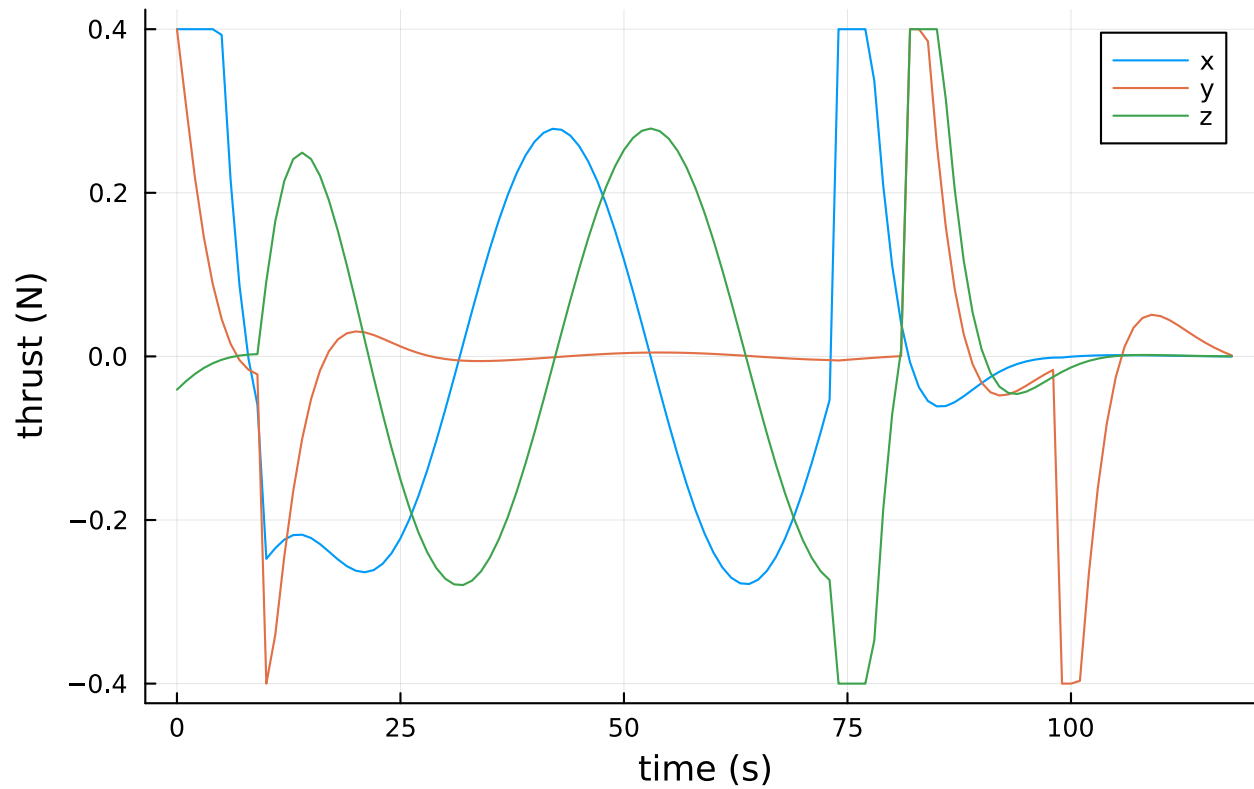
Positions (LQR)



Velocities (LQR)



Control (LQR)



[Info: Listening on: 127.0.0.1:8713, thread id: 1

└ Info: MeshCat server started. You can open the visualizer by visiting the following URL in your browser:

└ <http://127.0.0.1:8713>



Test Summary:	Pass	Total	Time
LQR rendezvous	6	6	0.5s

```
Out[ ]: Test.DefaultTestSet("LQR rendezvous", Any[], 6, false, false, true, 1.708630657357702e9, 1.708630657845608e9, false)
```

Part C: Convex Trajectory Optimization (15 pts)

Now we are going to assume that we have a perfect model (assume there is no sim to real gap), and that we have a perfect state estimate. With this, we are going to solve our control problem as a convex trajectory optimization problem.

$$\min_{x_{1:N}, u_{1:N-1}} \sum_{i=1}^{N-1} \left[\frac{1}{2} (x_i - x_{ref,i})^T Q (x_i - x_{ref,i}) + \frac{1}{2} u_i^T R u_i \right] (x_N - x_{ref,N}) \quad (7)$$

$$\text{st } x_1 = x_{IC} \quad (8)$$

$$x_{i+1} = Ax_i + Bu_i \quad \text{for } i = 1, 2, \dots, N-1 \quad (9)$$

$$u_{min} \leq u_i \leq u_{max} \quad \text{for } i = 1, 2, \dots, N-1 \quad (10)$$

$$x_i[2] \leq x_{goal}[2] \quad \text{for } i = 1, 2, \dots, N \quad (11)$$

$$x_N = x_{goal} \quad (12)$$

Where we have an LQR cost, an initial condition constraint ($x_1 = x_{IC}$), linear dynamics constraints ($x_{i+1} = Ax_i + Bu_i$), bound constraints on the control ($\leq u_i \leq u_{max}$), an ISS collision constraint ($x_i[2] \leq x_{goal}[2]$), and a terminal constraint ($x_N = x_{goal}$). This problem is convex and we will setup and solve this with `Convex.jl`.

```
In [ ]: """
Xcvx,Ucvx = convex_trajopt(A,B,X_ref,x0,xg,u_min,u_max,N)

setup and solve the above optimization problem, returning
the solutions X and U, after first converting them to
vectors of vectors with vec_from_mat(X.value)
"""
function convex_trajopt(A::Matrix, # discrete dynamics A
    B::Matrix, # discrete dynamics B
    X_ref::Vector{Vector{Float64}}, # reference trajectory
    x0::Vector, # initial condition
    xg::Vector, # goal state
    u_min::Vector, # lower bound on u
    u_max::Vector, # upper bound on u
    N::Int64, # length of trajectory
)::Tuple{Vector{Vector{Float64}}, Vector{Vector{Float64}}} # return Xcvx,Ucvx

    # get our sizes for state and control
    nx,nu = size(B)
```

```

@assert size(A) == (nx, nx)
@assert length(x0) == nx
@assert length(xg) == nx

# LQR cost
Q = diagm(ones(nx))
R = diagm(ones(nu))

# variables we are solving for
X = cvx.Variable(nx,N)
U = cvx.Variable(nu,N-1)

# TODO: implement cost
obj = 0
for i = 1:N-1
    obj += cvx.quadform(X[:,i] - X_ref[i], Q) + cvx.quadform(U[:,i], R)
end

# create problem with objective
prob = cvx.minimize(obj)

# TODO: add constraints with prob.constraints +=
prob.constraints += [X[:,1] == x0]
prob.constraints += [X[:,N] == xg]
for i = 1:N-1
    prob.constraints += [X[:,i+1] == A*X[:,i] + B*U[:,i]]
    for j = 1:nu
        prob.constraints += [u_min[j] <= U[j,i]]
        prob.constraints += [U[j,i] <= u_max[j]]
    end
    prob.constraints += [X[2,i] <= xg[2]]
end

cvx.solve!(prob, ECOS.Optimizer; silent_solver = true)

X = X.value
U = U.value

```

```

Xcvx = vec_from_mat(X)
Ucvx = vec_from_mat(U)

return Xcvx, Ucvx
end

@testset "convex trajopt" begin

    # create our discrete time model
    dt = 1.0
    A,B = create_dynamics(dt)

    # get our sizes for state and control
    nx,nu = size(B)

    # initial and goal states
    x0 = [-2;-4;2;0;0;.0]
    xg = [0,-.68,3.05,0,0,0]

    # bounds on U
    u_max = 0.4*ones(3)
    u_min = -u_max

    # problem size and reference trajectory
    N = 100
    t_vec = 0:dt:((N-1)*dt)
    X_ref = desired_trajectory(x0,xg,N,dt)

    # solve convex trajectory optimization problem
    X_cvx, U_cvx = convex_trajopt(A,B,X_ref, x0,xg,u_min,u_max,N)

    X_sim = [zeros(nx) for i = 1:N]
    X_sim[1] = x0
    for i = 1:N-1
        X_sim[i+1] = A*X_sim[i] + B*U_cvx[i]
    end

    # -----plotting/animation-----

```



```

Xm = mat_from_vec(X_sim)
Um = mat_from_vec(U_cvx)
display(plot(t_vec,Xm[1:3,:]',title = "Positions",
            xlabel = "time (s)", ylabel = "position (m)",
            label = ["x" "y" "z"]))
display(plot(t_vec,Xm[4:6,:]',title = "Velocities",
            xlabel = "time (s)", ylabel = "velocity (m/s)",
            label = ["x" "y" "z"]))
display(plot(t_vec[1:end-1],Um',title = "Control",
            xlabel = "time (s)", ylabel = "thrust (N)",
            label = ["x" "y" "z"]))

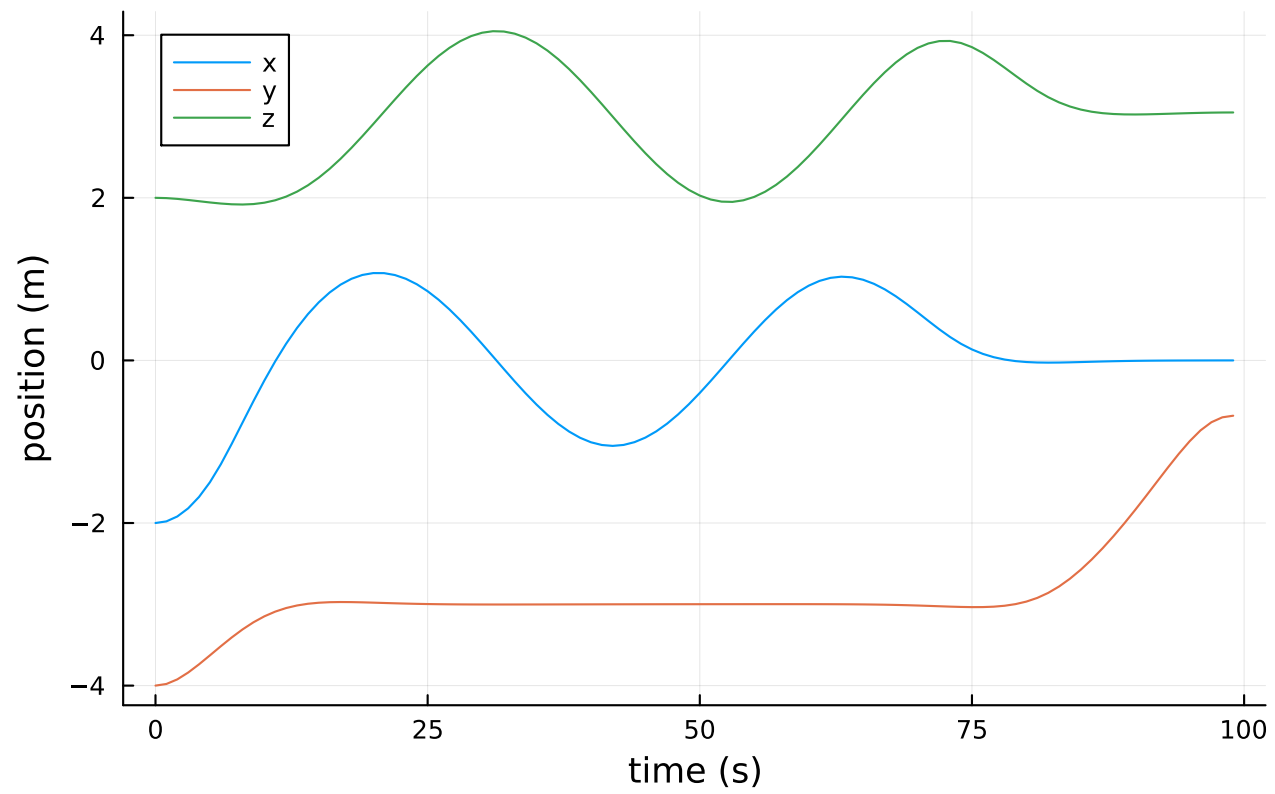
display(animate_rendezvous(X_sim, X_ref, dt;show_reference = false))
# -----plotting/animation-----

@test maximum(norm.( X_sim .- X_cvx, Inf)) < 1e-3
@test norm(X_sim[end] - xg) < 1e-3 # goal
xs=[x[1] for x in X_sim]
ys=[x[2] for x in X_sim]
zs=[x[3] for x in X_sim]
@test maximum(ys) <= (xg[2] + 1e-3)
@test maximum(zs) >= 4 # check to see if you did the circle
@test minimum(zs) <= 2 # check to see if you did the circle
@test maximum(xs) >= 1 # check to see if you did the circle
@test maximum(norm.(U_cvx,Inf)) <= 0.4 + 1e-3 # control constraints satisfied

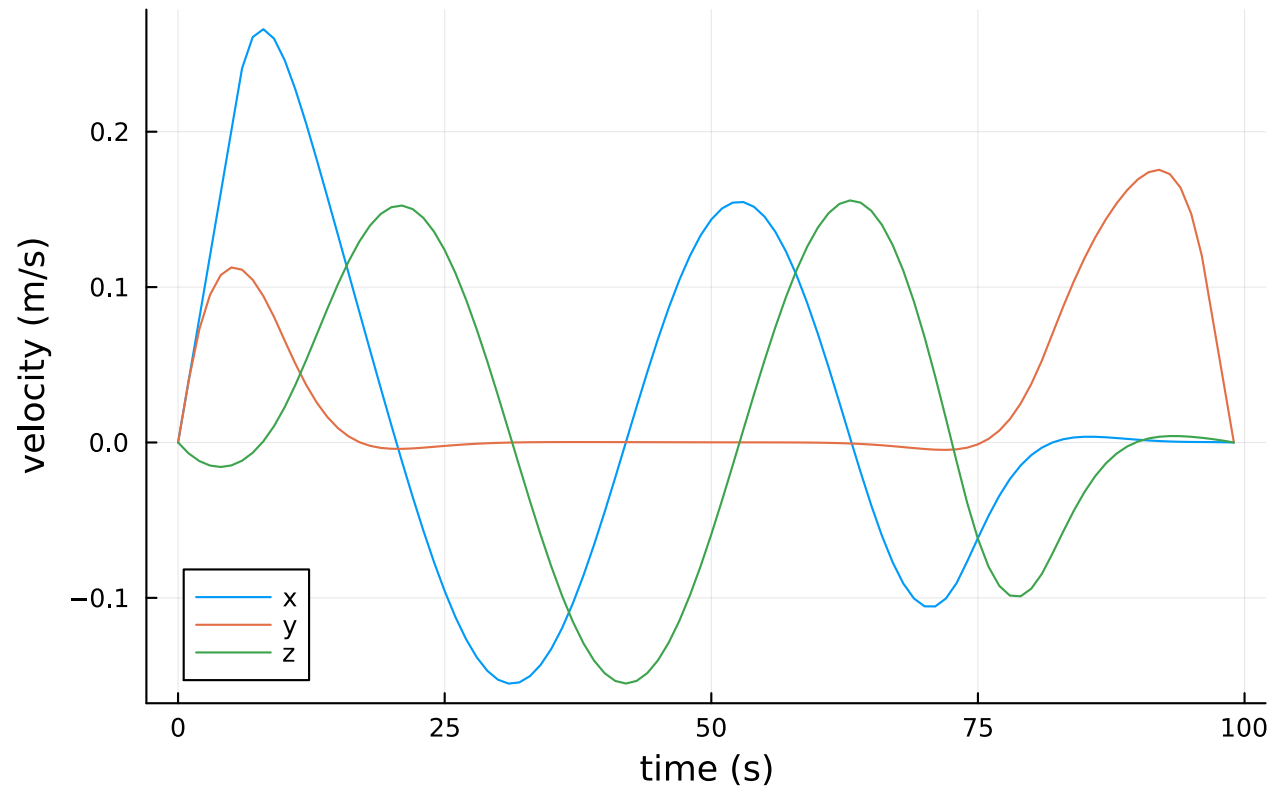
end

```

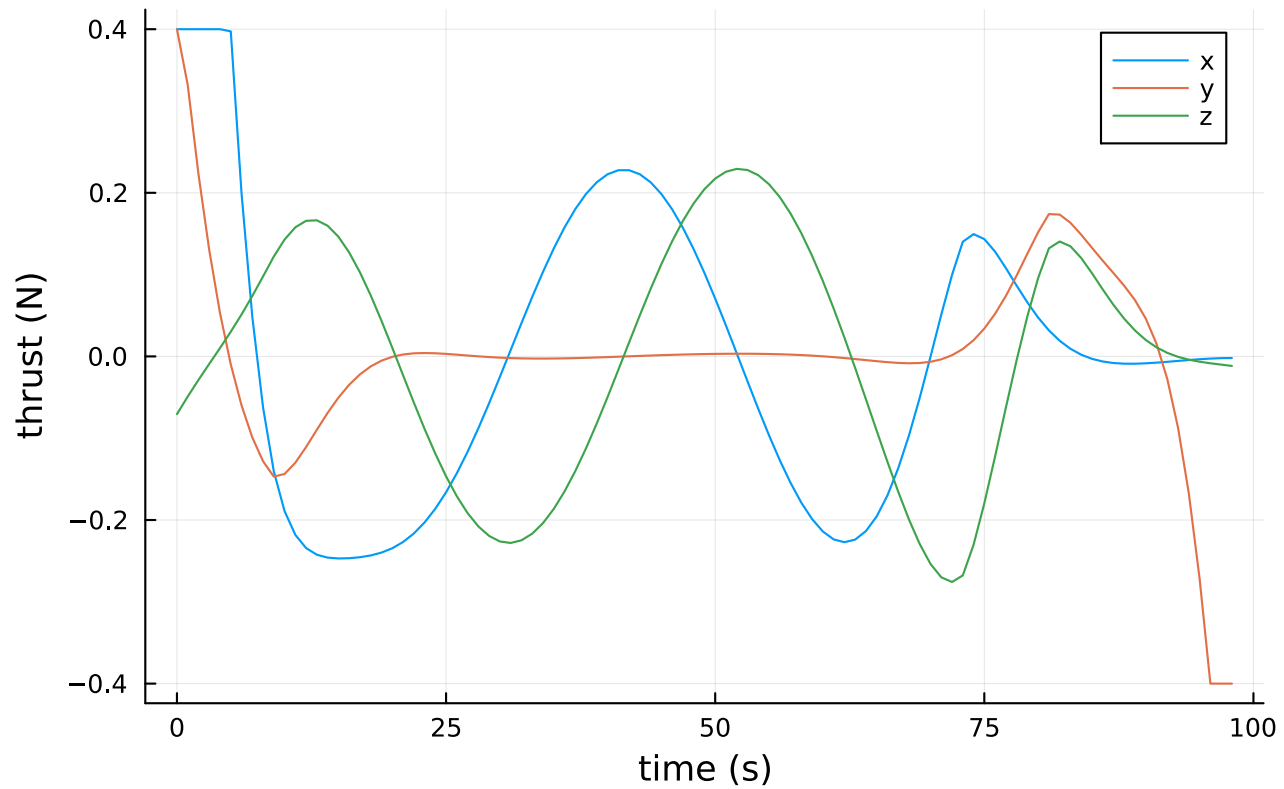
Positions



Velocities



Control



[Info: Listening on: 127.0.0.1:8714, thread id: 1

└ Info: MeshCat server started. You can open the visualizer by visiting the following URL in your browser:

└ <http://127.0.0.1:8714>



Test Summary:	Pass	Total	Time
convex trajopt	7	7	14.2s

```
Out[ ]: Test.DefaultTestSet("convex trajopt", Any[], 7, false, false, true, 1.708630663505171e9, 1.708630677714268e9, false)
```

Part D (5 pts): Short answer

1. List three reasons why an open loop policy wouldn't work well on a real system:

- Observation/control noise: The real system will have noise in the observations and control inputs, and an open loop policy will suffer from accumulating errors due to no feedback.
- Unmodeled dynamics: The real system will have unmodeled dynamics that the open loop policy will not be able to account for and will deviate from the desired trajectory.
- Parameter uncertainty: The real system might have different parameters than the model, and the open loop policy cannot compensate for this.

2. For convex trajectory optimization, give three examples of convex cost functions we can use:

- The L1 norm on the control inputs ($\|u\|_1$) and tracking error ($\|x - x_{goal}\|_1$)
- The L2 norm on the control inputs ($\|u\|_2$) and tracking error ($\|x - x_{goal}\|_2$)
- The log-sum-exp cost function ($\log(\sum_i e^{x_i})$)

1. List three things that convex trajectory optimization can do that LQR cannot:

- System with control limit: Convex trajectory optimization can handle control limits, while LQR cannot.
- System with state constraints: Convex trajectory optimization can handle state constraints, while LQR cannot.
- Highly nonlinear systems: Convex trajectory optimization can handle highly nonlinear systems, while LQR cannot.

4. Say we have the following convex trajectory optimization problem:

$$\min_{x_{1:N}, u_{1:N-1}} \sum_{i=1}^{N-1} \left[\frac{1}{2} (x_i - x_{ref,i})^T Q (x_i - x_{ref,i}) + \frac{1}{2} u_i^T R u_i \right] + \frac{1}{2} (x_N - x_{ref,N})^T Q_f (x_N - x_{ref,N}) \quad (13)$$

$$\text{st } x_1 = x_{IC} \quad (14)$$

$$x_{i+1} = Ax_i + Bu_i \quad \text{for } i = 1, 2, \dots, N-1 \quad (15)$$

$$x_{min} \leq x_i \leq x_{max} \quad \text{for } i = 1, 2, \dots, N \quad (16)$$

$$u_{min} \leq u_i \leq u_{max} \quad \text{for } i = 1, 2, \dots, N-1 \quad (17)$$

If the optimal solution to this problem does not violate any either the state or control bounds (the $x_{min} \leq x_i \leq x_{max}$ and $u_{min} \leq u_i \leq u_{max}$ constraints), how will it differ from the finite-horizon LQR solution?

A: The output of the convex trajectory optimization would be a open-loop policy that is feasible and optimal for the given horizon, while the finite-horizon LQR solution would be a feedback policy that is optimal for the given horizon.

Part E: Convex MPC (20 pts)

In part C, we solved for the optimal rendezvous trajectory using convex optimization, and verified it by simulating it in an open loop fashion (no feedback). This was made possible because we assumed that our linear dynamics were exact, and that we had a perfect estimate of our state. In reality, there are many issues that would prevent this open loop policy from being successful.

Together, these factors result in a "sim to real" gap between our simulated model, and the real model. Because there will always be a sim to real gap, we can't just execute open loop policies and expect them to be successful. What we can do, however, is use Model Predictive Control (MPC) that combines some of the ideas of feedback control with convex trajectory optimization.

A convex MPC controller will set up and solve a convex optimization problem at each time step that incorporates the current state estimate as an initial condition. For a trajectory tracking problem like this rendezvous, we want to track x_{ref} , but instead of optimizing over the whole trajectory, we will only consider a sliding window of size N_{mpc} (also called a horizon). If $N_{mpc} = 20$, this means our convex MPC controller is reasoning about the next 20 steps in the trajectory. This optimization problem at every timestep will start by taking the relevant reference trajectory at the current window from the current step i , to the end of the window $i + N_{mpc} - 1$. This slice of the reference trajectory that applies to the current MPC window will be called $\tilde{x}_{ref} = x_{ref}[i, (i + N_{mpc} - 1)]$.

$$\min_{x_{1:N}, u_{1:N-1}} \sum_{i=1}^{N-1} \left[\frac{1}{2} (x_i - \tilde{x}_{ref,i})^T Q (x_i - \tilde{x}_{ref,i}) + \frac{1}{2} u_i^T R u_i \right] + \frac{1}{2} (x_N - \tilde{x}_{ref,N})^T Q (x_N - \tilde{x}_{ref,N}) \quad (18)$$

$$\text{st } x_1 = x_{IC} \quad (19)$$

$$x_{i+1} = Ax_i + Bu_i \quad \text{for } i = 1, 2, \dots, N-1 \quad (20)$$

$$u_{min} \leq u_i \leq u_{max} \quad \text{for } i = 1, 2, \dots, N-1 \quad (21)$$

$$x_i[2] \leq x_{goal}[2] \quad \text{for } i = 1, 2, \dots, N \quad (22)$$

where N in this case is N_{mpc} . This allows for the MPC controller to "think" about the future states in a way that the LQR controller cannot. By updating the reference trajectory window (\tilde{x}_{ref}) at each step and updating the initial condition (x_{IC}), the MPC controller is able to "react" and compensate for the sim to real gap.

You will now implement a function `convex_mpc` where you setup and solve this optimization problem at every timestep, and simply return u_1 from the solution.

```
In [ ]: """
`u = convex_mpc(A,B,X_ref_window,xic,xg,u_min,u_max,N_mpc)`

setup and solve the above optimization problem, returning the
first control  $u_1$  from the solution (should be a length nu
Vector{Float64}).
"""
function convex_mpc(A::Matrix, # discrete dynamics matrix A
                   B::Matrix, # discrete dynamics matrix B
                   X_ref_window::Vector{Vector{Float64}}, # reference trajectory for this window
                   xic::Vector, # current state x
                   xg::Vector, # goal state
                   u_min::Vector, # lower bound on u
                   u_max::Vector, # upper bound on u
                   N_mpc::Int64, # length of MPC window (horizon)
                   )::Vector{Float64} # return the first control command of the solved policy

    # get our sizes for state and control
    nx,nu = size(B)

    # check sizes
    @assert size(A) == (nx, nx)
    @assert length(xic) == nx
    @assert length(xg) == nx
    @assert length(X_ref_window) == N_mpc

    # LQR cost
    Q = diagm(ones(nx))
    R = diagm(ones(nu))

    # variables we are solving for
    X = cvx.Variable(nx,N_mpc)
    U = cvx.Variable(nu,N_mpc-1)

    # TODO: implement cost function
```



```

obj = 0
for i = 1:N_mpc-1
    obj += cvx.quadform(X[:,i] - X_ref_window[i], Q) + cvx.quadform(U[:,i], R)
end

# create problem with objective
prob = cvx.minimize(obj)

# TODO: add constraints with prob.constraints +=
prob.constraints += [X[:,1] == xic]
prob.constraints += [X[:,N_mpc] == xg]
for i = 1:N_mpc-1
    prob.constraints += [X[:,i+1] == A*X[:,i] + B*U[:,i]]
    for j = 1:nu
        prob.constraints += [u_min[j] <= U[j,i]]
        prob.constraints += [U[j,i] <= u_max[j]]
    end
    prob.constraints += [X[2,i] <= xg[2]]
end

# solve problem
cvx.solve!(prob, ECOS.Optimizer; silent_solver = true)

# get X and U solutions
X = X.value
U = U.value

# return first control U
return U[:,1]
end

@testset "convex mpc" begin

    # create our discrete time model
    dt = 1.0
    A,B = create_dynamics(dt)

```

```

# get our sizes for state and control
nx,nu = size(B)

# initial and goal states
x0 = [-2;-4;2;0;0;.0]
xg = [0,-.68,3.05,0,0,0]

# bounds on U
u_max = 0.4*ones(3)
u_min = -u_max

# problem size and reference trajectory
N = 100
t_vec = 0:dt:((N-1)*dt)
X_ref = [desired_trajectory(x0,xg,N,dt)...,[xg for i = 1:N]...]

# MPC window size
N_mpc = 20

# sim size and setup
N_sim = N + 20
t_vec = 0:dt:((N_sim-1)*dt)
X_sim = [zeros(nx) for i = 1:N_sim]
X_sim[1] = x0
U_sim = [zeros(nu) for i = 1:N_sim-1]

# simulate
@showprogress "simulating" for i = 1:N_sim-1

    # get state estimate
    xi_estimate = state_estimate(X_sim[i], xg)

    # TODO: given a window of N_mpc timesteps, get current reference trajectory
    X_ref_tilde = X_ref[i:i+N_mpc-1]

    # TODO: call convex mpc controller with state estimate
    u_mpc = convex_mpc(A,B,X_ref_tilde,xi_estimate,xg,u_min,u_max,N_mpc)

    # commanded control goes into thruster model where it gets modified

```

```

    U_sim[i] = thruster_model(X_sim[i], xg, u_mpc)

    # simulate one step
    X_sim[i+1] = A*X_sim[i] + B*U_sim[i]
end


# -----plotting/animation-----
Xm = mat_from_vec(X_sim)
Um = mat_from_vec(U_sim)
display(plot(t_vec,Xm[1:3,:]',title = "Positions",
            xlabel = "time (s)", ylabel = "position (m)",
            label = ["x" "y" "z"]))
display(plot(t_vec,Xm[4:6,:]',title = "Velocities",
            xlabel = "time (s)", ylabel = "velocity (m/s)",
            label = ["x" "y" "z"]))
display(plot(t_vec[1:end-1],Um',title = "Control",
            xlabel = "time (s)", ylabel = "thrust (N)",
            label = ["x" "y" "z"]))

display(animate_rendezvous(X_sim, X_ref, dt;show_reference = false))
# -----plotting/animation-----

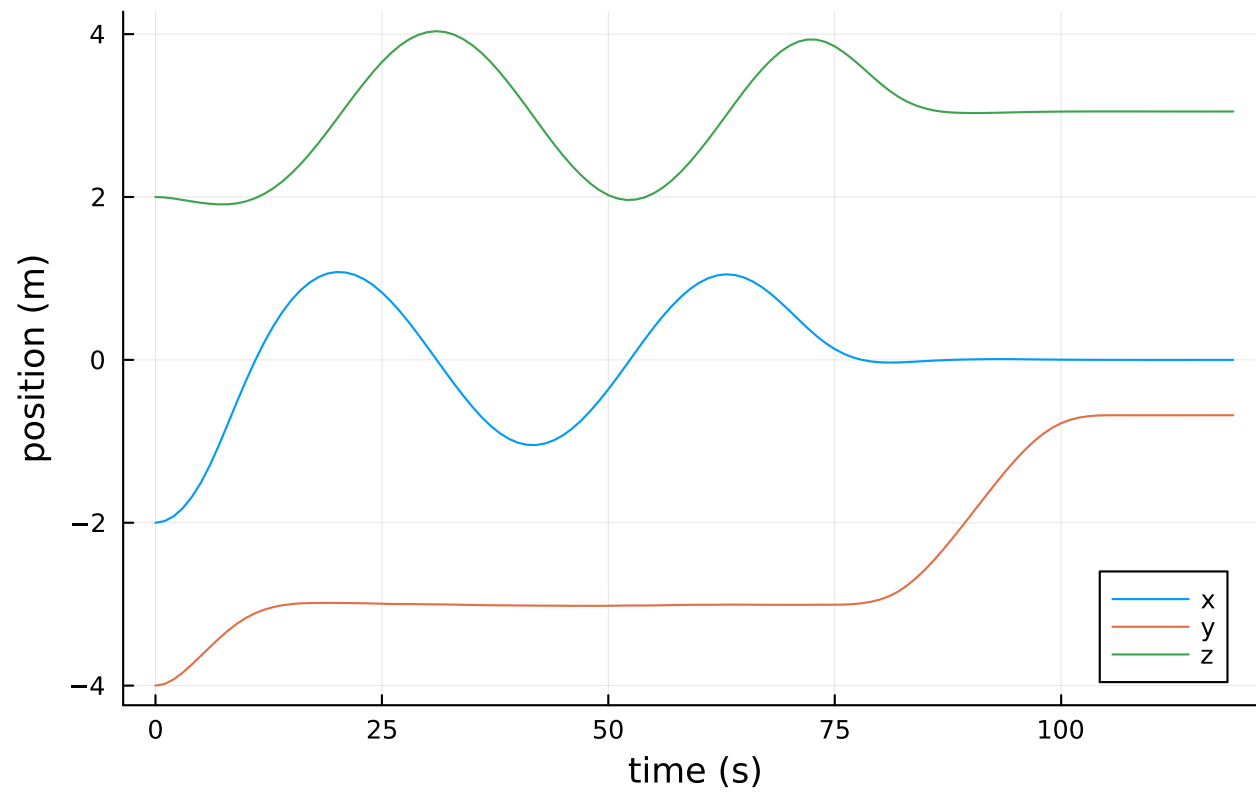
# tests
@test norm(X_sim[end] - xg) < 1e-3 # goal
xs=[x[1] for x in X_sim]
ys=[x[2] for x in X_sim]
zs=[x[3] for x in X_sim]
@test maximum(ys) <= (xg[2] + 1e-3)
@test maximum(zs) >= 4 # check to see if you did the circle
@test minimum(zs) <= 2 # check to see if you did the circle
@test maximum(xs) >= 1 # check to see if you did the circle
@test maximum(norm.(U_sim,Inf)) <= 0.4 + 1e-3 # control constraints satisfied

end

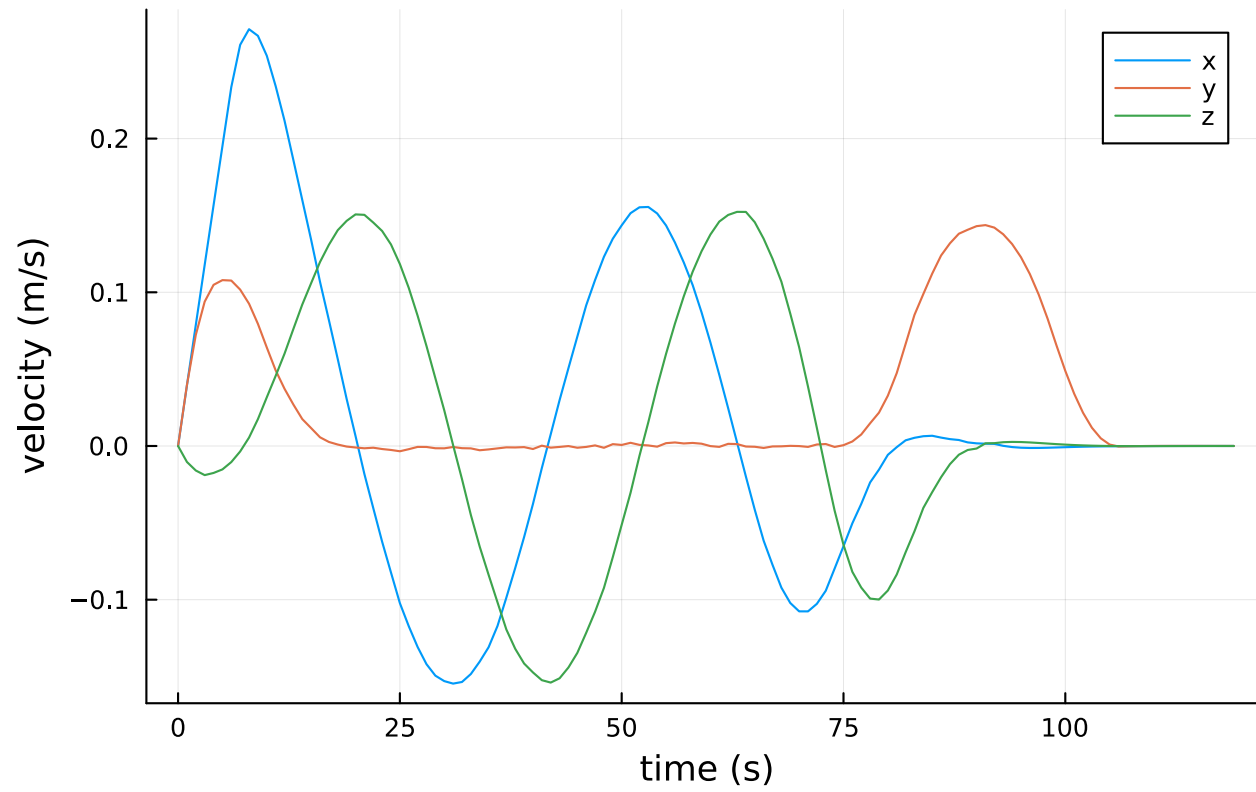
```

simulating 100%|  | Time: 0:00:02

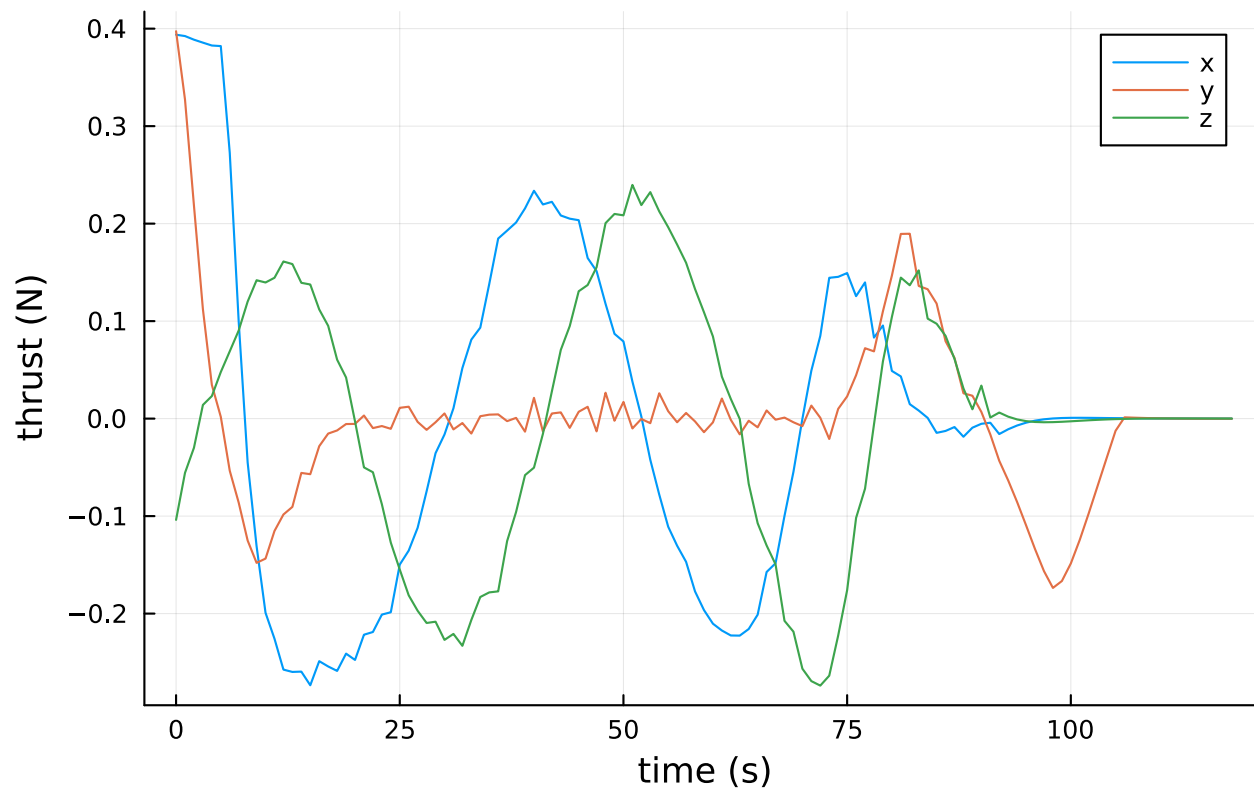
Positions



Velocities



Control



[Info: Listening on: 127.0.0.1:8715, thread id: 1

└ Info: MeshCat server started. You can open the visualizer by visiting the following URL in your browser:

└ <http://127.0.0.1:8715>

Test Summary: | **Pass** **Total** **Time**
convex mpc | 6 6 2.9s

Out[]: Test.DefaultTestSet("convex mpc", Any[], 6, false, false, true, 1.70863067795458e9, 1.708630680835651e9, false)